

Style Transfer

CS 20: TensorFlow for Deep Learning Research Lecture 9 2/9/2017

Announcements

Assignment 2 is out. It's fun, but tricky. Start early.

Sign up for check-ins/IGs with the course staff! cs20-win1718-staff@lists.stanford.edu

Guest lectures next week



Alec Radford OpenAI Topic: GANs 2/9



Danijar Hafner Google Brain Topic: Variational Autoencoder 2/14

Agenda

TFRecord

Getting to know each other!

Style Transfer





TFRecord

What's TFRecord

- 1. The recommended format for TensorFlow
- 2. Binary file format

What's TFRecord

- 1. The recommended format for TensorFlow
- 2. Binary file format a serialized tf.train.Example protobuf object

Why binary

• make better use of disk cache

Why binary

- make better use of disk cache
- faster to move around

Why binary

- make better use of disk cache
- faster to move around
- can handle data of different types e.g. you can put both images and labels in one place

• Feature: an image

• Label: a number

```
# Step 1: create a writer to write tfrecord to that file
writer = tf.python io.TFRecordWriter(out file)
# Step 2: get serialized shape and values of the image
shape, binary image = get image binary(image file)
# Step 3: create a tf.train.Features object
features = tf.train.Features(feature={'label': int64 feature(label),
                                    'shape': bytes feature(shape),
                                    'image': bytes feature(binary image)})
# Step 4: create a sample containing of features defined above
sample = tf.train.Example(features=features)
# Step 5: write the sample to the tfrecord file
writer.write(sample.SerializeToString())
writer.close()
```

```
# Step 1: create a writer to write tfrecord to that file
writer = tf.python io.TFRecordWriter(out file)
# Step 2: get serialized shape and values of the image
shape, binary image = get image binary(image file)
# Step 3: create a tf.train.Features object
features = tf.train.Features(feature={'label': int64 feature(label),
                                    'shape': bytes_feature(shape),
                                    'image': bytes feature(binary image)})
# Step 4: create a sample containing of features defined above
sample = tf.train.Example(features=features)
# Step 5: write the sample to the tfrecord file
writer.write(sample.SerializeToString())
writer.close()
```

Serialize different data type into byte strings

```
def _int64_feature(value):
    return tf.train.Feature(int64_list=tf.train.Int64List(value=[value]))

def _bytes_feature(value):
    return tf.train.Feature(bytes_list=tf.train.BytesList(value=[value]))
```

Read TFRecord

Using TFRecordDataset

Read TFRecord

```
dataset = tf.data.TFRecordDataset(tfrecord_files)
dataset = dataset.map(_parse_function)
```

Parse each tfrecord_file into different features that we want

In this case, a tuple of (label, shape, image)

Read TFRecord

See o8_tfrecord_example.py



Assignment 2: Style Transfer

Bringing Impressionism to Life with Neural Style Transfer in Come Swim

Bhautik J Joshi*

Research Engineer, Adobe

Director, Come Swim

David Shapiro
Producer, Starlight Studios







Figure 1: Usage of Neural Style Transfer in Come Swim; left: content image, middle: style image, right: upsampled result. Images used with permission, (c) 2017 Starlight Studios LLC & Kristen Stewart.

Yes, that Kristen Stewart!



Deadpool



Guernica



Deadpool and Guernica

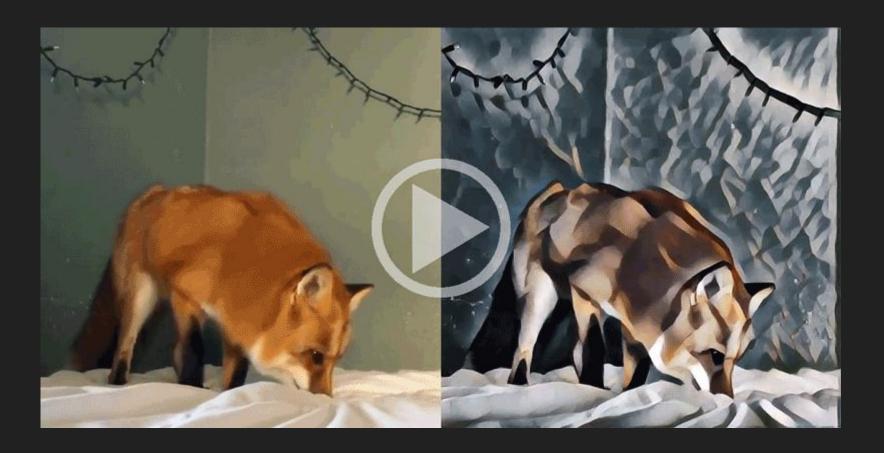












Style Transfer

The math is aight but the implementation is tricky

Mathy stuff

Find a new image:

- whose content is closest to the content image and
- whose style is closest to the style image

It's all about the loss functions

Content loss

Measure the content loss between the content of the generated image and the content of the content image

Style loss

Measure the style loss between the style of the generated image and the style of the style image

WHAT'S CONTENT?

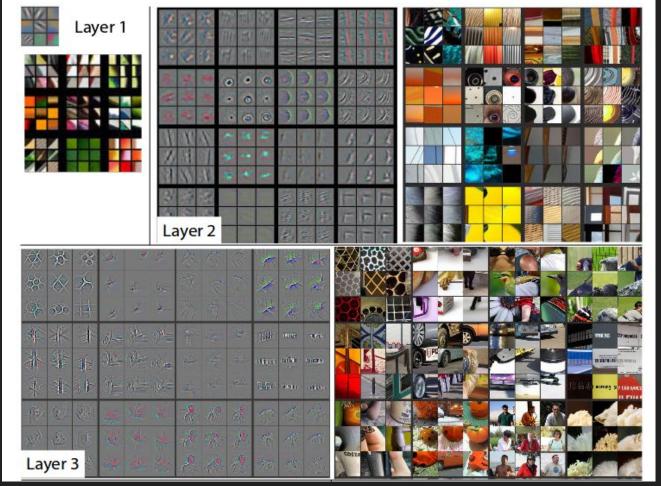


30



Feature maps

A convolutional network has many layers, each layer is a function that extracts certain features



Content/style of an image

Feature visualization have shown that:

- lower layers extract features related to content
- higher layers extract features related to style

Loss functions revisited

- Content loss
 - Measure the loss between the feature maps in the content layer of the generated image and the content image
- Style loss
 - Measure the loss between the feature maps in the style layers of the generated image and the style image

Loss functions revisited

Content loss

To measure the content loss between **the feature map in the content layer** of the generated image and the content image

Paper: 'conv4_4'

• Style loss

To measure the style loss between **the gram matrices of feature maps in the style layers** of the generated image and the style image

Paper: ['conv1_1', 'conv2_1', 'conv3_1', 'conv4_1' and 'conv5_1']

Loss functions revisited

Content loss

To measure the content loss between **the feature map in the content layer** of the generated image and the content image

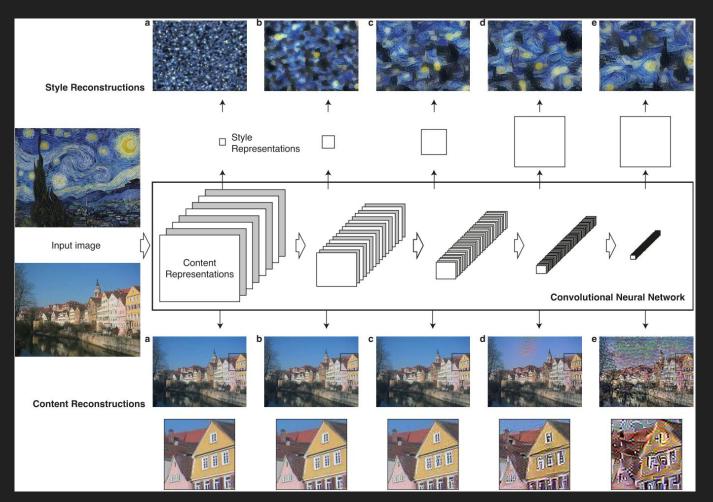
Paper: 'conv4_4'

• Style loss

To measure the style loss between **the gram matrices of feature maps in the style layers** of the generated image and the style image

Paper: ['conv1_1', 'conv2_1', 'conv3_1', 'conv4_1' and 'conv5_1']

Weighted sum. Give more weight to deeper layers E.g. 1.0 for 'conv1_1', 2.0 for 'conv2_1', ...



How to find these magic feature maps?

Use pretrained weights (functions) such as VGG, AlexNet, GoogleNet

Loss functions revisited

Content loss

$$\mathcal{L}_{content}(ec{p},ec{x},l) = rac{1}{2} \sum_{i,j} \left(F_{ij}^l - P_{ij}^l
ight)^2$$

• Style loss

$$E_{l} = rac{1}{4N_{l}^{2}M_{l}^{2}}\sum_{i,j}\left(G_{ij}^{l}-A_{ij}^{l}
ight)^{2}$$

$$\mathcal{L}_{style}(ec{a},ec{x}) = \sum_{l=0}^{L} w_l E_l$$

Optimizer

Optimizes the initial image to minimize the combination of the two losses

$$\mathcal{L}_{total}(\vec{p}, \vec{a}, \vec{x}) = \alpha \mathcal{L}_{content}(\vec{p}, \vec{x}) + \beta \mathcal{L}_{style}(\vec{a}, \vec{x})$$

Do not optimize the weights!

Tricky implementation details

1. Train input instead of weights

Tricky implementation details

- 1. Train input instead of weights
- 2. Multiple tensors share the same variable to avoid assembling identical subgraphs

Tricky implementation details

- 1. Train input instead of weights
- 2. Multiple tensors share the same variable to avoid assembling identical subgraphs
- 3. Use pre-trained weights (from VGG-19)
 - a. Weights and biases already loaded for you
 - b. They are numpy, so need to be converted to tensors
 - c. Must not be trainable!!

Progress



Cool story, bro. So what?

Fun applications

- Snapchat filters
- Google photos
- Movies!!!

Is art exclusively a human domain?

Next class

GANs by Alec Radford!

Feedback: chiphuyen@cs.stanford.edu

Thanks!